Errata for "Cerebro: A Data System for Optimized Deep Learning Model Selection"

Supun Nakandala, Yuhao Zhang, and Arun Kumar
University of California, San Diego
{snakanda, yuz870, arunkk}@eng.ucsd.edu

We discovered that there was an inconsistency in the communication cost formulation for the decentralized fine-grained training method in **Table 2** of our paper [1]. We used Horovod as the archetype for decentralized fine-grained approaches, and its correct communication cost is higher than what we had reported. So, we amend the communication cost of decentralized fine-grained to $\mathbf{2km(p-1)|S|\left\lceil\frac{|D|}{bp}\right\rceil}$, instead of $kmp|S|\left\lceil\frac{|D|}{bp}\right\rceil$.

With this correction, **Table 2** of our paper should be corrected as follows, which uses the same notation.

Table 1: Communication cost analysis of MOP and other approaches. *Full replication. †Remote reads. ‡Parameters for the example: $k = 20$, $|S| = 20$, $p = 10$, $m = 1$GB, $\langle D\rangle = 1$TB, and $|D|/b = 100$K.

	Comm. Cost	Example‡
Model Hopper Parallelism	$kmp\|S\| + m\|S\|$	4 TB
Task Parallelism (FR*)	$p\langle D\rangle + m\|S\|$	10 TB
Task Parallelism (RR†)	$k\|S\|\langle D\rangle + m\|S\|$	400 TB
Bulk Synchronous Parallelism	$2kmp\|S\|$	8 TB
Centralized Fine-grained	$2kmp\|S\|\left\lceil\frac{\|D\|}{bp}\right\rceil$	80 PB
Decentralized Fine-grained	$\mathbf{2km(p-1)\|S\|\left\lceil\frac{\|D\|}{bp}\right\rceil}$	**72 PB**

Also, the last two paragraphs of Section 2 that refers to the above table should be corrected as follows:

> All PS-style approaches have *high communication* due to their centralized all-to-one communications, which is proportional to the number of mini-batches and orders of magnitude higher than BSP, e.g., **10,000x** in Table 2.
>
> Decentralized Fine-grained. The best example is Horovod. It adopts HPC-style techniques to enable synchronous all-reduce SGD. While this approach is bandwidth optimal, communication latency is still proportional to the number of workers, and the synchronization barrier can become a bottleneck. The total communication overhead is also proportional to the number of mini-batches and orders of magnitude higher than BSP, e.g., **9,000x** in Table 2.

The above amendments are purely in the conceptual exposition and do not affect any technical findings, empirical results, or conclusions in the paper.

Proceedings of the VLDB Endowment, Vol. 13, No. 11
ISSN 2150-8097.
DOI: https://doi.org/10.14778/3407790.3407816

1. REFERENCES

[1] S. Nakandala, Y. Zhang, and A. Kumar. Cerebro: A data system for optimized deep learning model selection. *Proc. VLDB Endow.*, 13(12):21592173, July 2020.