

Lotan: Bridging the Gap between GNNs and Scalable Graph Analytics Engines

Yuhao Zhang

University of California, San Diego
yuz870@eng.ucsd.edu

Arun Kumar

University of California, San Diego
akk018@ucsd.edu

ABSTRACT

Recent advances in Graph Neural Networks (GNNs) have changed the landscape of modern graph analytics. The complexity of GNN training and the scalability challenges have also sparked interest from the systems community, with efforts to build systems that provide higher efficiency and schemes to reduce costs. However, we observe that many such systems basically “reinvent the wheel” of much work done in the database world on scalable graph analytics engines. Further, they often tightly couple the scalability treatments of graph data processing with that of GNN training, resulting in entangled complex problems and systems that often do not scale well on one of those axes.

In this paper, we ask a fundamental question: How far can we push existing systems for scalable graph analytics and deep learning (DL) instead of building custom GNN systems? Are compromises inevitable on scalability and/or runtimes? We propose Lotan, the first scalable and optimized data system for full-batch GNN training with *decoupled scaling* that bridges the hitherto siloed worlds of graph analytics systems and DL systems. Lotan offers a series of technical innovations, including re-imagining GNN training as query plan-like dataflows, execution plan rewriting, optimized data movement between systems, a GNN-centric graph partitioning scheme, and the first known GNN model batching scheme. We prototyped Lotan on top of GraphX and PyTorch. An empirical evaluation using several real-world benchmark GNN workloads reveals a promising nuanced picture: Lotan significantly surpasses the scalability of state-of-the-art custom GNN systems, while often matching or being only slightly behind on time-to-accuracy metrics in some cases. We also show the impact of our system optimizations. Overall, our work shows that the GNN world can indeed benefit from building on top of scalable graph analytics engines. Lotan’s new level of scalability can also empower new ML-oriented research on ever-larger graphs and GNNs.

PVLDB Reference Format:

Yuhao Zhang and Arun Kumar. Lotan: Bridging the Gap between GNNs and Scalable Graph Analytics Engines. PVLDB, 16(11): xxxx - xxxx, 2023. doi:xxxx/xxxx

PVLDB Artifact Availability:

The source code, data, and/or other artifacts have been made available at <https://github.com/makemebitter/lotan>.

This work is licensed under the Creative Commons BY-NC-ND 4.0 International License. Visit <https://creativecommons.org/licenses/by-nc-nd/4.0/> to view a copy of this license. For any use beyond those covered by this license, obtain permission by emailing info@vldb.org. Copyright is held by the owner/author(s). Publication rights licensed to the VLDB Endowment.

Proceedings of the VLDB Endowment, Vol. 16, No. 11 ISSN 2150-8097.
doi:xxxx/xxxx

1 INTRODUCTION

Graph Neural Networks (GNNs) have drastically shifted the landscape of advanced graph analytics. GNNs can provide powerful learned representations for graphs. In about a decade, GNNs have dominated many graph analytics leaderboards [21] for tasks ranging from lower-level ones, such as node classification and edge prediction, to graph-level tasks like graph classification or even graph generation. Applications span from video analytics [23], recommender systems [64, 70], drug discovery [34] and pandemic data analysis [67], to even crime prediction [56] with spatial-temporal graphs. Interest in GNNs is rising rapidly in many domains where data are naturally represented as graphs, such as social networks and molecular structures.

However, GNN models are tricky to scale [18, 63, 66], because of the sheer amount of computation and the immense memory pressure they exert on GPUs. A plethora of GNN systems was proposed to tackle these challenges [12, 24, 35, 39, 50, 60, 62, 75]. They express GNN workloads primarily as advanced matrix multiplications and rely on GPUs for execution. When GPU memory is insufficient to host the entire matrices and the intermediate results, one either resorts to distributed processing [24, 75] and/or spilling techniques [24, 60] that load/offload data from GPU accordingly.

What makes GNN training so hard to scale, and why do we need these dedicated systems for GNNs? First, graph data are irregularly shaped and non-IID, differentiating them from regular IID data modalities such as text and images, for which the state-of-the-art DL frameworks were designed. To tackle the data scalability issues, most DL frameworks employ distributed data-parallel schemes [1, 55]. However, data parallelism does not directly apply to graph data: graph partitions are not independent, and the training process involves cross-partition communications, depending on the input graph structure. Second, neural network backpropagation requires caching intermediate data during forward propagation. Depending on the graph data, these intermediates could be huge in size. Unlike models such as CNNs or Transformers designed for IID data, where input shape is often normalized and uniform, GNNs are highly input-dependent. They are tough to accommodate, as workloads are highly versatile and vary significantly in scale. Third, the “neighborhood explosion” and the over-smoothing problems [7, 8, 41] are also tricky to bypass; data dependency grows exponentially as the number of GNN layers grows, posing challenges to both scalability and efficiency.

In this paper, we make a critical observation that **many of the GNN’s challenges are, in fact, challenges of managing, moving, and handling the underlying graph data**. Nonetheless, existing custom GNN systems mix and couple the graph data and DL challenges. We observe several shortcomings of this worldview: first, many of these systems “reinvent the wheel” of much work

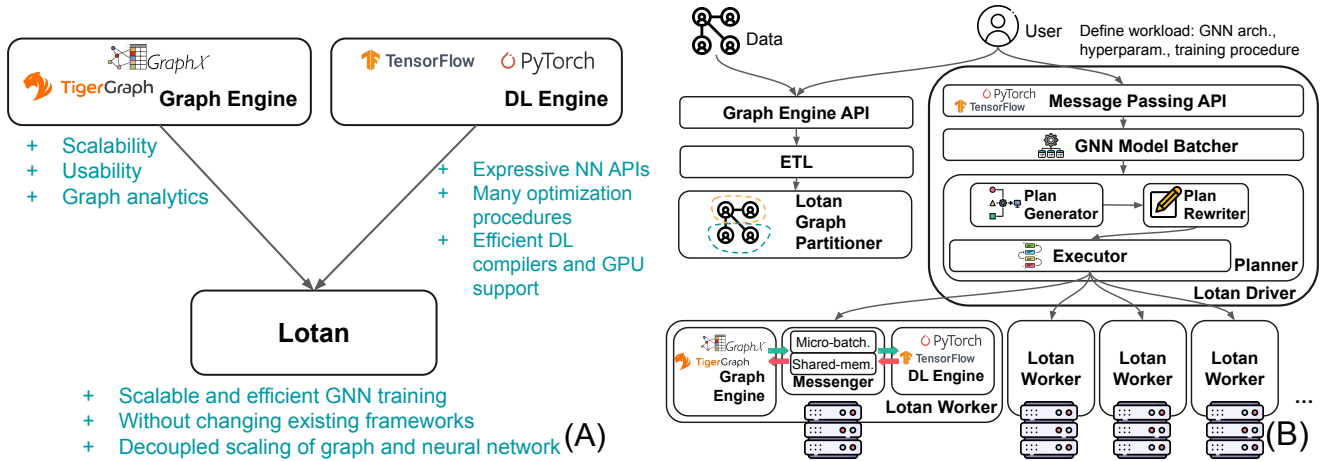


Figure 1: (A) Lotan bridges the gap between graph systems and DL systems. (B) The architecture of Lotan.

done in the database world on scalable graph analytics engines. Second, they often tightly couple the scalability treatments of graph data processing with that of GNN training, resulting in entangled, complex problems and systems that often do not scale well on one of those axes. GNN workloads, though drastically different from regular DNN workloads in data access patterns, are not too far away from non-NN graph analytics such as PageRank. As pointed out by prior work [39], most of the popular GNNs can be expressed under extended versions of graph programming models such as Gather-Apply-Scatter (GAS). Scaling “shallow” graph analytics is not a new topic: many graph data systems were designed for that purpose [15–17, 40]. However, to our knowledge, none of these systems provide general GNN support, nor do they handle DL operations, which, nowadays, are better reserved for frameworks such as TensorFlow and PyTorch. It would be prohibitively labor-intensive to build systems of their generality and performance from scratch. Furthermore, implementing native NN support within graph systems would lead to a similar problem of reinventing the wheel of DL systems research. Therefore, both stacks of software are needed: graph systems for graph challenges and DL systems for DL challenges. As Figure 1(A) explains, our work aims to bridge this gap.

System Desiderata. We envision a scalable GNN system with the following desiderata: (1) *Decoupled Scaling*: Scale graph and neural network parts by reusing existing industrial-strength tools. (2) *Usability*: Retain the ease of specification APIs of both graph and DL tools. (3) *Non-disruptive Integration*: No changes to the internal code of those tools. (4) *Speed and Accuracy*: Fast runtimes without sacrificing DL accuracy.

In this paper, we seek to answer a fundamental systems question: *How far can we go by pushing existing systems’ limits without modifications?* We propose a novel information system architecture for scalable GNN training with the **decoupling of graph and neural network**. Much like the famous **decoupling of compute and storage** in cloud computing, this decoupling enables us to tackle each side individually and allows them to scale independently. We carefully pick apart the graph and neural network dataflows in GNN training and re-imagine them as a “query plan” in our new

intermediate-level *global operator graph*. We dispatch the execution plan to an existing graph analytics engine and a DL framework without modifying their internal code. We built a distributed prototype system we call Lotan.

Overview of Lotan. Figure 1(B) illustrates our system architecture. The user interacts with Lotan through the APIs to specify their GNN workload. Our Planner then compiles it into graph and neural network operations and dispatches them to their separate execution Engines, which are existing graph and DL systems. Our Messenger handles the coordination and communications between the Engines. This way, Lotan can preserve all functionalities of the execution Engines, especially the graph data management functionalities such as graph manipulation, partitioning, and other non-NN graph analysis methods.

Additionally, Lotan provides a series of system optimizations to increase the runtime performance. The most important two are GNN-centric graph partitioning and GNN model batching.

GNN-centric graph partitioning. Distributed graph processing naturally comes with the problem of graph partitioning, which can dramatically affect the efficiency as sub-optimal partitioning leads to a huge amount of network communications. To this end, we propose a GNN-centric graph partitioning scheme and the corresponding Reverse Graph Propagation execution scheme for GNN training. Our method works by taking account of the asymmetry in data size between forward- and backward propagation. We will describe it in detail in Section 5.1.

GNN model batching. GNNs, like other DL methods, require extensive hyperparameter tuning, which involves training multiple models on the same dataset. These models overlap extensively in their data access patterns, and opportunities exist because data access is quite costly for GNNs. We propose the first GNN Model Batching technique to improve GPU utilization and reduce runtime for GNN model selection workloads. As far as we know, Lotan is the first system to optimize for the GNN model selection/hyperparameter tuning workloads and the first to explore model batching for GNNs. We will introduce it in Section 5.2.

Overall, this paper makes the following technical contributions:

- To the best of our knowledge, this is the first work to bridge the gap between existing graph data systems and DL systems and the first to formally decouple the scaling of graph and neural networks in GNN training. Lotan expands design freedom for GNN researchers and practitioners.
- We re-imagine large-scale GNN training from a data management standpoint and unpack the dataflows into a “query plan” representation. We then devise novel query rewriting and optimization techniques to improve scalability and efficiency.
- We propose one of the first GNN-centric graph partitioning schemes to reduce graph node replication and communications during GNN training.
- Furthermore, Lotan is the first GNN system to treat model selection workloads holistically and explore model batching techniques to improve training throughput.
- We perform an extensive evaluation to compare Lotan with prior industrial-strength systems and study the impact of our optimizations. The empirical results validate our system’s higher scalability and competitive time-to-accuracy performance on multiple workloads.

2 BACKGROUND

2.1 Graph Neural Networks

Graph Neural Networks (GNNs) are neural networks on graph data. In a nutshell, a GNN always tries to summarize the graph structure and/or the graph properties into a compact numerical representation called embeddings. GNNs can be categorized into spectral-based and spatial-based methods [66]. Spectral-based methods have roots in graph signal processing and rely on the graph Laplacian and Fourier transform for generating embeddings. Spatial-based methods are typically the applications of neural networks such as RNN, CNN, and GAN on graph data, with modifications to account for graph structure. The spatial-based methods are the more popular of the two categories [66] and will be the main focus of our system.

It is important to note that a GNN model can be ultimately expressed as a combination of graph processing (in the form of a modified Gather-Apply-Scatter programming model [39]) and DL operations. This is the basis of how our system attacks the problem; we compile a GNN training task into a global operator graph composed of graph operators and neural network operators and use existing systems for execution. More details on these concepts are in Section 3.2 and Section 4.

2.2 Distributed Graph Processing

GNN workloads are still a form of graph processing/analytics because they resemble many classical problems and share very similar data access patterns. To tackle the many similar challenges, non-GNN graph data systems [4, 9, 15, 17, 38, 54, 59] rely on distributed processing, and a critical problem is graph data partitioning.

There are two major graph partitioning schemes: edge-cut and vertex-cut. It is beyond the scope of this paper to fully cover the entire landscape of graph partitioning, so we only introduce the bare minimum background before we propose our own GNN-centric




Figure 2: Two graph partitioning schemes.

graph partitioning scheme in Section 5.1. Interested readers are directed to other literature on graph partitioning [5].

Edge-cut. Edge-cut partitioning affixes the location of vertices, and the edges at partitioning boundaries are replicated (or need to be remotely fetched when needed). Figure 2(A) illustrates this scheme. In Gather-Apply-Scatter workloads, the messages generated at vertices are sent across the partitioning boundary, resulting in cross-partition communications.

Vertex-cut. Vertex-cut partitioning is the alternative to edge-cut; it focuses on the edges and fixes their locations. As a trade-off, the vertices at the boundaries need to be replicated or at least remotely fetched when needed.

Note that a certain amount of cross-partition communication or data replication is inevitable, depending on the quality of the graph partitioning algorithm and the characteristics of the underlying graph. Graphs typically have much more edges than vertices in the real world. Therefore vertex-cut partitioning, which avoids edge replications, sometimes are more favorable in practice [17, 38, 53] for non-GNN workloads. Although a plethora of graph partitioning algorithms exists [5], they are seldom designed for GNN workloads. As a result, there is room for improvement. We will dive deep into the characteristics of GNN workloads in Section 5.1 and subsequently propose our graph partitioning scheme on top of the vertex-cut scheme.

2.3 GNN Training Systems

Plenty of systems have been proposed to tackle the efficiency and scalability challenges brought by GNNs. Generally speaking, there are two main camps within GNN system research: first is the scalability camp, which aims to tackle the scalability issues of full-batch GNNs [24, 58, 76]; they are usually distributed systems and focus on providing the capability to run GNN workloads that fail on other systems. Second is the efficiency camp, which mainly focuses on runtime speed and usually does not address scalability issues; they

Table 1: Comparison with prior art on key capabilities.

	License	GPU	Distributed	Sampling	Memory Hierarchy
Lotan	Open	✓	✓	Full	Disk-aware
DGL/DistDGL [62]	Open	✓	✓	Both	GPU-only
AliGraph/graph-learn [76]	Open	✓	✓	Mini-batch	GPU-only
PSGraph [25]	N/A	✓	✓	Mini-batch	GPU-only
GraphScope [68]	Open	✓	✓	Mini-batch	GPU-only
Sancus [50]	Open	✓	✓	Full	GPU-only
PipeGCN [61]	Open	✓	✓	Full	GPU-only
Dorylus [58]	Open	✗	✓ (Serverless)	Full	N/A
ROC [24]	Open	✓	✓	Full	DRAM-aware
P ³ [13]	N/A	✓	✓	Mini-batch	GPU-only
DeepGalois [20]	N/A	✗	✓	Full	DRAM-only
Pytorch Geometric [12]	Open	✓	✗	Both	GPU-only
NeuGraph [39]	N/A	✓	✗	Full	DRAM-aware
PaGraph [3, 35]	Open	✓	✗	Mini-batch	DRAM-aware
MariusGNN [60]	Open	✓	✗	Mini-batch	Disk-aware

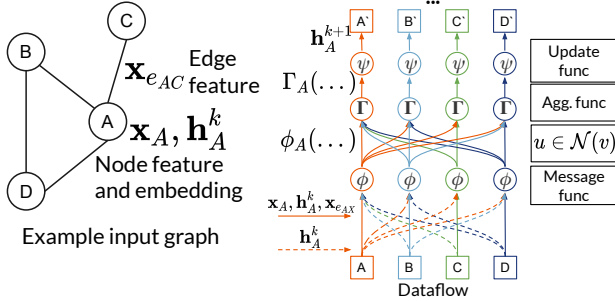


Figure 3: (A) An example input graph to a spatial-based GNN. (B) Dataflow diagram of a message passing GNN.

often assume that the entire workload can comfortably fit in GPU memory/main memory [35, 39, 42, 50].

We summarize the comparisons between these systems in Table 1, and we will discuss them in more detail in Section 8. We evaluate these systems by several axes: (1) License, whether the system is open-source and usable for tests. (2) GPU, whether the system has GPU support. (3) Distributed, whether the system support distributed processing. (4) Sampling, whether the system targets full-batch or mini-batch GNN training. (5) Memory Hierarchy, whether the system is secondary storage aware. We leave the performance tests and numbers to Section 7. In this paper, we argue that many of these systems are reinventing the wheel with custom-built graph data systems and are still facing scalability issues with larger datasets or models. Lotan is more closely related to the scalability camp, but it differs from the prior art in technical contributions and our architecture design, which can utilize existing, established systems.

3 GNN APIS AND PROGRAMMING MODEL

3.1 GNN Interface

The first system issue we need to address is *how do we express a GNN model in a standardized way?* One of the most common and general abstractions is known as the Message Passing interface [14]. It is also widely adopted in GNN system literature and the de-facto standard. The Message Passing interface defines a GNN using *update rule*, an equation that tells us how to update a graph node’s embedding:

$$\mathbf{h}_v^k = \psi(\mathbf{x}_v^k, \prod_{u \in \mathcal{N}(v)} \phi(\mathbf{h}_v^{k-1}, \mathbf{h}_u^{k-1}, \mathbf{x}_{e_{vu}})), \quad (1)$$

where ψ, ϕ, Γ are potentially learnable and differentiable functions, Γ is further required to be commutative and associative. ψ is called the update function, ϕ the message function and Γ the aggregate function. Note Equation 1 covers primarily spatial convolutional GNNs, on which we focus. Some other forms of GNNs, such as spectral-based methods [66], cannot be easily and efficiently expressed in the same framework. We leave the question of how to support those GNNs in future work.

Equation 1 is the interface we expose to the user. They will need to define the three functions ψ, ϕ, Γ using APIs and operators provided by the system to be introduced next. Depending on the nature of these functions, our system can do plan rewriting and optimizations to boost performance. More details are in Section 4.2. Figure 3(A) shows an example input graph and Figure 3(B) shows the conceptual dataflow of a GNN being learned on it.

Batched Message Passing. In practice, we find it much more beneficial to rewrite Equation 1 in a batched and vectorized format, especially for better utilization of GPU:

$$\mathbf{H}_v^k = \Psi(\mathbf{X}_v^k, \prod^* \Phi(\mathbf{H}_v^{k-1}, \mathbf{H}_u^{k-1}, \mathbf{X}_{e_{vu}})), \quad (2)$$

where $\mathbf{H}_v^k, \mathbf{X}_v^k, \mathbf{H}_v^{k-1}, \mathbf{H}_u^{k-1}$, and $\mathbf{X}_{e_{vu}}$ are all matrices which are batched forms of their corresponding vectors, they have shape $B \times _$, where B is the batch size and $_$ is the dimension of the respective vectors. Ψ, Φ, Γ^* are the batched (vectorized) form of functions ψ, ϕ, Γ .

3.2 Lotan’s Internal Programming Model

It is not obvious how to parallelize Equation 1, mainly due to the neighborhood aggregation steps that are only native to graph processing systems. Some prior work re-cast the equation into bulk linear algebra [12, 62, 75, 76]. However, they often encounter huge scalability issues due to the potentially colossal graph size and the resulting sizes of matrix multiplications. On the other hand, it is not unfamiliar to see existing GNN systems using Gather-Apply-Scatter (GAS) abstraction or its extension [39] to translate Equation 1 into an executable plan. However, none of these abstractions capture the potentially costly data transfer operations between the graph and DL execution Engines that Lotan relies on. We need to account for these operations correctly so that we can better understand the problem. Toward this goal, we propose a new programming model that roots in the decoupling of compute and storage and the decoupling of graph and NN. Our abstraction involves three main stages: (1) Graph propagation with Scatter-Gather-Collect, (2) DL propagation with ApplyEdge-Aggregation-ApplyVertex, and (3) Pipe and Join. To implement these operators, we build them upon existing operators in the Graph and DL Engines.

Scatter-Gather-Collect. During this stage, the Graph Engine does a regular Scatter and Gather as in GAS for the graph propagation portion of a GNN. Instead of Apply in GAS, here it is followed by a Collect operation, where the Graph Engine, depending on the

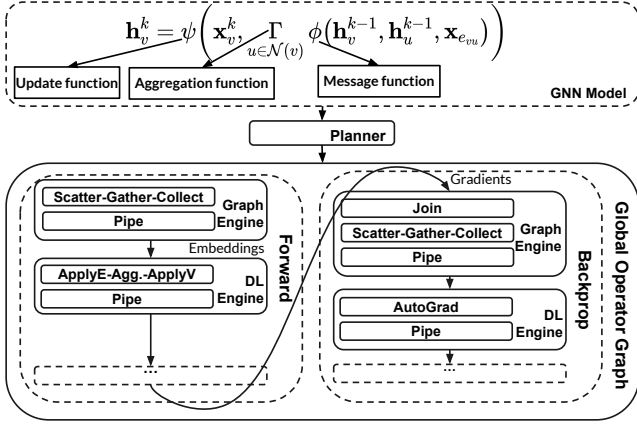


Figure 4: Global operator graph of end-to-end GNN training.

specifications of the GNN, collects and packs relevant data to hand over to the DL Engine. This arises because GNN operations are placed on two different Engines in our system. The DL Engine operations (such as aggregation) have data dependencies up to the Graph Engine to resolve. The Graph Engine then needs to collect all the data from each graph node and their neighboring nodes and sends them to the DL Engine. Conceptually, this stage is primarily for implementing the neighborhood scope $u \in \mathcal{N}(v)$ in Equation 1.

ApplyEdge-Aggregation-ApplyVertex. In this stage, the DL Engine receives data from the Graph Engine and applies the GNN functions on the data. ApplyEdge implements the per-edge function ϕ ; similarly, ApplyVertex implements the per-vertex function ψ . Aggregation implements the neighborhood aggregation function Γ .

Pipe and Join. We need operations for data transfer at the Graph and DL Engine’s boundary. From the Graph Engine to DL Engine, we need a Pipe operation that, as the name suggests, pipes data to the DL Engine and the results back to the Graph Engine. Then within the Graph Engine, a Join operation is needed to incorporate the data, as the order of data may not be preserved during the Pipe. We will cover these in detail in Section 4.3.

One important note is that this separation of stages is not fixed; some operations can be eliminated, some can be re-ordered, and some can be pushed down. We will explore all these opportunities for optimizations in Section 4.2.

3.3 Global Operator Graph and Execution

With all the introduced abstractions, we can now compile an entire GNN training workload into a global operator graph with the operators mentioned above. A Planner, to be discussed in Section 4.2, will generate this graph from the user input expressed in the GNN message-passing interface.

Figure 4 shows the full operator graph for end-to-end GNN training, using the operators defined in Section 3.2. Data (embeddings during the forward-propagation, and gradients during the back-propagation) is sent back and forth between the Graph Engine and DL Engine. The Graph Engine is in charge of the graph aggregation by running Gather-Scatter-Apply under the hood for both forward-

and back-propagation and collects all the necessary data for the DL Engine to consume, represented by the Collect operator. During the forward-propagation, the DL Engine handles the ApplyEdge, Aggregation, and ApplyVertex functions and subsequently does back-propagation with their AutoGrad capabilities. Both Engines run independently and are unaware of each other. They can run on the same set of machines, and the operators are parallelized independently. To coordinate the Engines and to provide a bridge for data transfer, we build a Messenger component for our system, to be introduced in Section 4.3.

4 SYSTEM ARCHITECTURE

Lotan has 3 main components: (1) External Engines, which are existing graph processing systems and DL frameworks **without modifications**. (2) Planner, where Lotan creates and optimizes the execution plan of a GNN training workload. (3) Messenger, where Lotan reconciles the Graph Engine and the DL Engine and facilitates efficient data transmission between them.

4.1 External Engines

These engines are what Lotan relies on and improves on for tackling many scalability challenges of GNN training. We only use these engines’ public interfaces and treat them as black boxes. This way, we use them without modifications and drastically increase the portability and generality of Lotan while preserving all the features provided by both Engines.

Graph Engine. The Graph Engine is an external graph data system that Lotan relies on for graph-related operations and scalability challenges. It can be a graph processing system or a graph DBMS, as long as it provides public interfaces for (1) Gather-Apply-Scatter (GAS) operators. (2) Operations that export data to external systems. Additionally, it should provide scalable solutions for large-scale graph analytics. Often, such Engines conveniently provide various data system features such as data partitioning and distribution, fault tolerance, memory management, and disk spilling. Most of today’s graph analytics systems/graph DBMS meet these criteria. Examples include Spark’s GraphX [17], Giraph [15], TigerGraph [9], and Neo4j [49]. We choose GraphX for our prototype because, first, it is open-source software and has an active user community. Second, it is easy to use and piggybacks on the popular and familiar Spark ecosystem. Our approach is general and easily applicable to other graph analytics engines.

Deep Learning Engine. To handle the challenge from the neural network part of a GNN, we adopt an external DL system/framework. We use this system for forward propagation activation computing and back-propagation gradient computing with their autograd capabilities. By using an existing DL framework, we automatically make available the rich DL libraries and GPU support such a framework comes with. Furthermore, these systems can offer an out-of-box solution for distributed model training via their data-parallel capabilities [32, 55]. TensorFlow and PyTorch are both prominent examples of such systems and both are applicable. We pick PyTorch due to its dominant popularity in the GNN community.

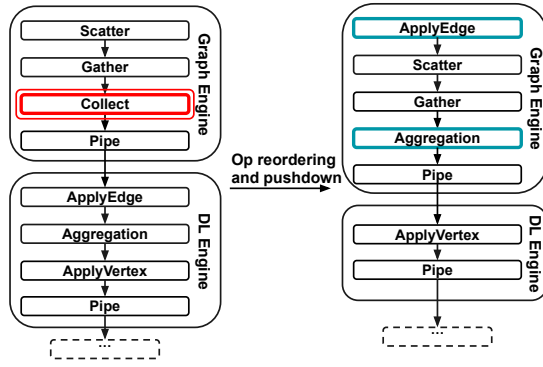


Figure 5: An example of plan rewrites. Note the Collect operator is rewritten with ApplyEdge and Aggregation.

4.2 Planner

At the heart of Lotan is the Planner, inspired by query planners in database research. Close to the concept of a DBMS query optimizer/planner, we need to weigh the potential query plans and choose the optimal one. The general idea is to assign relative costs for each stage of the execution and then determine the final cost estimate. However, in this case, the plan search space is much more limited, and it is favorable to do operator pushdowns whenever possible. Therefore, we find simple heuristics sufficient and no sophisticated cost estimation is needed. To complete the study, we still try to model the costs, but primarily for curiosity and a deeper understanding of the problem. We will also verify some of the observations from our cost models with experiments in our technical report [72].

Plan Generation and Rewrites. Plan generation is usually trivial, as GNN training comprises mostly sequential stages, as Figure 4 shows. Opportunities for optimization exist; depending on the nature of the GNN model, the execution plan can be rewritten. We only consider the two most obvious cases of plan rewriting: operator reordering and pushdown.

Equation 1 gives the most general definition of a GNN, and Figure 3(B) is the most stringent ordering of the three functions from the Message function to the Aggregation function to the Update function. However, because all these functions can be neural networks, they can only be handled by the DL Engine. Both the Message and Aggregation functions require neighbor information; we will also have to collect all the edges, features, and embeddings in the Graph Engine and ship them to the DL Engine. For this general case, our operator graph writes as Figure 4. This is an expensive plan due to the Collect operator and the size of data movement between the two Engines. However, if the Message and Aggregation functions are both unparameterized and therefore do not require training, we can push down these functions to the Graph Engine and drastically save costs. Figure 5 illustrates this scheme. We will test plan rewrites with experiments in Section 7.2.1 and see that it contributes to substantial performance gains.

Cost Estimation. To calculate the costs of a plan, we first evaluate the costs of individual stages respectively, then aggregate them

together. A stage is defined as the sub-operator graph between two boundaries of data movement. The costs are estimated using: (1) the data graph’s information, such as the number of nodes and vertices and the average degree. (2) specifications of the GNN, such as the number of layers and the number of parameters involved in the neural network. (3) The DRAM and GPU RAM limit, network/disk bandwidth, and the number of concurrent CPU threads available (degree of parallelism). Due to space constraints, we will highlight the main observations in Section 6, but leave the details to our technical report [72]. Within the cost models, a few factors are situation-dependent and, therefore, cannot be very well estimated. One can resort to logs of past runs of the same model and graph for more accurate costs.

4.3 Micro-batch Processing and Messenger

One critical question of utilizing existing systems is how to reconcile them; each comes with its input/output interfaces, data formats, memory layouts, and other specifications. Further, the DL Engine heavily favors batched data input for higher utilization and throughput, while the data comes off the Graph Engine as streams to reduce memory footprint. This means we must convert the data stream to and from data batches. We also need to keep the order of data consistent during both the forward pass and backward pass stages.

To our best knowledge, this is the first time the data movement issues between graph data systems and DL systems are being studied. We adopt and synthesize existing techniques and optimizations to solve the novel problems mentioned above. We build a component called Messenger. We apply a series of system optimizations to the Messenger: It uses non-blocking, async sockets and shared memory to communicate with the DL Engine for overlapping computation with communication and to reduce throttling. The details of this component can be found in our technical report [72].

5 SYSTEM OPTIMIZATIONS

5.1 GNN-centric Graph Partitioning and Reverse Graph Backpropagation

Graph partitioning is a vital part of distributed graph processing, as it dramatically impacts the volume of data replication and communications. Most existing graph partitioning schemes are not designed with GNNs in mind, resulting in suboptimal performance. We propose a novel graph partitioning and training execution scheme for GNNs, named Reverse Graph Backpropagation (RGB). This technique applies to vertex-cut-based Graph Engines [16, 17] such as GraphX, which we use for prototyping. Our method is based on two key observations: first, neural network training consists of a forward- and a back-propagation phase; the two phases have inverted dataflow. Second, during GNN training, graph node data are updated, and the data size changes between phases, which leads to an asymmetry in replication costs.

We start from the well-accepted hash-based 1D edge partitioning [53], where we hash partition all the nodes and then colocate all edges based on their sources. Figure 6 illustrates the strategy. During forward propagation, each node’s property and messages it sends are its node embeddings, which are 1-dimensional vectors. No node replication happens, but two cross-partition messages take

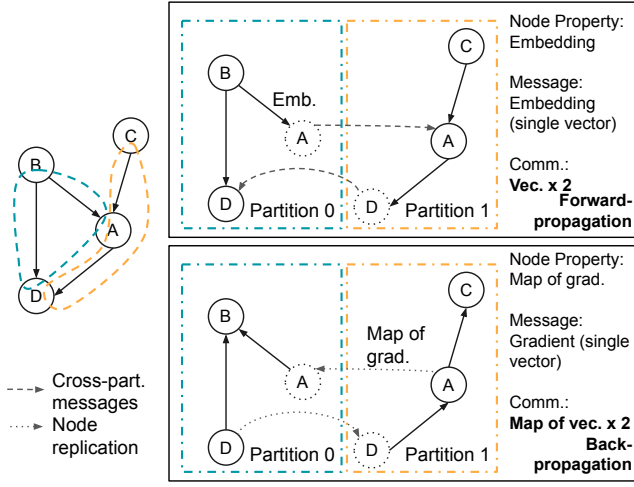


Figure 6: Regular 1D source hash partitioning and dataflow.

place. During the back-propagation, the data flow is inverted. However, each node updates its properties to gradients returned from the DL Engine (such updates happen in-partition and do not incur cross-partition communications). These gradients are hash maps of vectors and, compared to the embeddings, are d (node degree) times larger. For a realistic graph, the average degree can easily be around 100. Because dataflow is inverted and the partitioning is not, heavy cross-partition communications would occur.

To address this performance issue caused by the asymmetry, we propose our novel GNN-centric Graph Partitioning scheme and the way to backpropagate through it, described as follows:

- (1) Create a reverse graph (each edge reversed) from the original graph.
- (2) Do a regular hash partitioning on the reverse graph: first, hash partition all the nodes and place them; second, partition the edges based on their sources so that all edges originating from the same source colocate in the same partition.
- (3) Finally, we partition the original graph's edges in the same manner but keep the node partitions generated from the reverse graph.
- (4) We run the forward propagation as usual on the original graph. However, we run the back-propagation on the reverse graph.

This way, there is drastically reduced communication during back-propagation, where the most significant bottleneck could arise. Depending on the circumstances, communication costs might increase for forward propagation but are offset by back-propagation savings. We keep the node placements consistent between phases, otherwise extra cross-partition communication will occur. Figure 7 illustrates our approach. Regarding cross-partition communications, we only have single vectors instead of hashmaps of vectors. The example shows a directed graph, but the same logic still applies to undirected graphs.

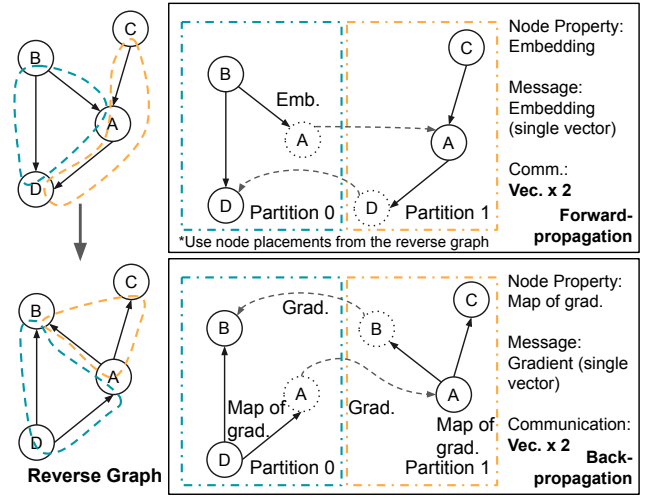


Figure 7: GNN-centric Graph Partitioning and dataflow.

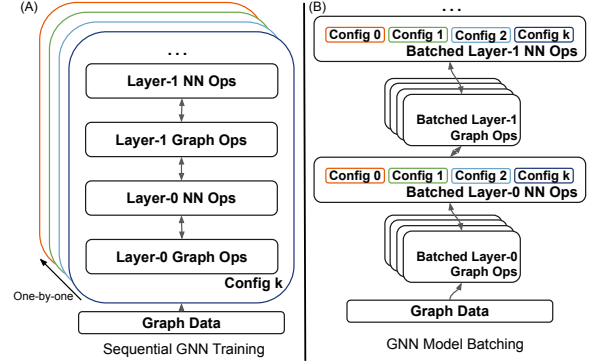


Figure 8: (A) Sequential training (B) Model Batching.

5.2 GNN Model Batching

GNNs, like any other neural network, rely on careful and extensive hyperparameter tuning for the best accuracy performance. Consequently, the workloads are often multiple-model explorations. Each model has its different set of neural network hyperparameters. When running hyperparameter tuning workloads, existing systems take a sequential approach: training them one-by-one. Figure 8(A) shows it. There is wasted potential for improvements: First, models in a hyperparameter search workload share identical data access patterns, and re-using these routines can amortize the overheads. Second, many GNN workloads have relatively low neural network components, often leaving the GPU underutilized. For DL methods on IID and Euclidean data, many systems [46, 73] have been developed to optimize for model selection workloads. However, these techniques do not apply as they assume IID data.

To address these issues, we propose GNN Model Batching. Model batching [48] is a technique to increase GPU utilization for IID models. To our best knowledge, we are the first to explore the same possibility for GNNs. We devise a model batching scheme to combine the models within a hyperparameter search workload.




Figure 9: Learning curves for the chosen model on the test set. (A) ogbn-products-GCN. (B) ogbn-products-GIN. (C) ogbn-arxiv-GCN. (D) ogbn-arxiv-GIN. Corresponding learning curves on the validation set are presented in our technical report [72].

Figure 8(B) depicts it. We run multiple models simultaneously on the model-batched version of the regular graph and NN operators. All data transmitted between the Graph Engine and DL Engine are also batched together. The models can then share all the data access operations to amortize costs.

6 ANALYSIS OF COST MODELS

To better understand the problem, we model the various costs of GNN training: replication, computational, memory, and overheads. As mentioned earlier in Section 4.2, these models are not used in the Planner and only for deeper understanding and further experiment evaluation. Due to space constraints, we leave the tedious details and equations to our technical report [72]. We present a summary of two key observations about our cost model that we will validate empirically later.

Effect of Number of Partitions. The number of data partitions interplays with system performance in two ways: First, for large-scale dataset, more partitions are required to reduce memory pressure. Second, increasing the number of partitions will also increase the degree of parallelism and utilization because our Graph Engine uses one thread per partition.

To put it into an equation. We have the total computational cost for an execution plan with partitions:

$$W_P = \frac{W}{P} \max\left(\frac{P}{ML}, 1\right) + f_{\text{overhead}}\left(\frac{P}{ML}\right), \quad (3)$$

where W is the total amount of work (unit: time), $\frac{W}{P}$ is each partition’s amount of work, $\frac{P}{ML}$ is the total amount of tasks each machine gets, $\max\left(\frac{P}{ML}, 1\right)$ is the total amount of rounds each machine executes. Without losing generality, assume f_{overhead} follows a monotonic increase along with P . We can then reason that as the number of partitions P increases; the overall runtime would first decrease and then increase.

We see precisely this behavior in our tests. Due to space constraints, the experiment details are moved to our technical report [72]. Due to the intertwined effects of this one parameter, the runtime behavior becomes non-linear and difficult to capture with simple cost models. Instead, we use rule-based heuristics to tune the number of partitions: we set it to be the same as the total number of CPU cores of the entire cluster unless more partitions

are required to alleviate the memory pressure. Fortunately, GNN workloads are highly predictable in runtime and resource consumption. If necessary, one can always do test runs (for 1 or 2 epochs of training is more than sufficient) to figure out the optimal config setting.

Effect of Model Batching. The intermediate embedding/gradient size of a GNN model greatly impacts runtime performance. Because of GNN Model Batching, Lotan can have inflated intermediates sizes. To be precise, the intermediate sizes will be multiplied by the model batching size. Consequently, for model batching, we expect to see a scaling up when increasing model batching size due to higher utilization until the returns diminish due to overheads. We run experiments and show the results in Section 7.2.3 and will see the expected behavior.

7 EXPERIMENTS AND EVALUATION

Prior Art. Out of the distributed GNN training systems discussed in Section 2.3, we show comparisons to the SOTA: DistDGL [75], AliGraph [76], and Sancus [50]. We excluded all systems that do not support distributed training and those without public release. Despite the best effort, we could not set up and use ROC [24], with a similar situation reported in [58]. Sancus [50] and PipeGCN [61] should be comparable systems, both with approximated processing, while others (including ours) are with exact processing, and we pick the former for benchmarking. Note that both DistDGL and AliGraph are primarily mini-batch GNN systems. Although DistDGL can run full-batch training, it fails almost all our workloads. Therefore, we use it with the mini-batch setting. Mini-batch training is mathematically different from full-batch training. But we put in our best effort for a fair comparison by tuning the mini-batch size to their advantage and using standard benchmark metrics agnostic of the training scheme.

Datasets. We use three of the standard benchmarking datasets from OGB [21], which has become the go-to place for graph datasets for benchmarking. We use ogbn-products, ogbn-arxiv, and ogbn-papers100M. Additionally, we also include datasets reddit [18] and amazon [19], the original amazon dataset is not shipped in graph form, and we converted it to graph after acquiring a recipe from the authors of [24, 58]. The prior art also commonly use these datasets

Table 2: End-to-end test results. TLE: time limit exceeded (48 hrs per model). *For these tests, all models within the workload learned too slowly and got terminated too soon. To not understate the system’s capability, we make an exception and include a separate run for a fixed 500 epochs and then pick the best valid accuracy checkpoint. †These tests would take an unreasonable amount of time to finish. Therefore, we did not train them to converge and only reported the throughput numbers.

Dataset Summary	Dataset	Model	System	Test Acc. (%)	Total Runtime (hr)	Throughput (epoch/hr)	CPU Util. (%)	GPU Util. (%)	Disk R/W (GB/hr)	Network (GB/hr)
#Nodes: 169.3K #Edges: 1.1M Avg. Degree: 13.7	ogbn-arxiv	GCN	Lotan	69.28	1.90	924.85	18.47	5.79	4294.66	4203.44
			DistDGL	68.49	0.33	1912.50	21.39	15.80	3.45	5959.28
			AliGraph	68.60	171.44	1.59	5.35	6.73	3.48	44.46
			*Sancus	*55.23	*0.79	1855.67	5.98	89.97	3.00	19706.78
		GIN	Lotan	71.22	3.71	557.38	18.11	5.93	3123.74	4959.36
			DistDGL	43.64	0.15	1035.97	20.13	16.65	4.13	6109.54
			*DistDGL	*69.26	*2.76	-	-	-	-	-
			AliGraph	Fail	-	-	-	-	-	-
#Nodes: 232.9K #Edges: 114.6M Avg. Degree: 492.9	reddit	GCN	Lotan	94.50	16.81	77.82	30.80	0.85	6173.45	4154.67
			DistDGL	Fail	-	-	-	-	-	-
			AliGraph	Fail	-	-	-	-	-	-
			Sancus	92.67	0.04	1408.69	5.89	75.38	113.57	14160.26
		GIN	Lotan	94.91	23.15	50.16	30.47	0.83	6366.86	4083.77
			DistDGL	Fail	-	-	-	-	-	-
			AliGraph	Fail	-	-	-	-	-	-
			Fail	-	-	-	-	-	-	-
#Nodes: 2.4M #Edges: 61.8M Avg. Degree: 50.5	ogbn-products	GCN	Lotan	75.59	63.82	16.22	49.58	1.75	6779.44	4989.52
			DistDGL	75.32	365.53	0.34	8.01	27.94	3.6	4548.72
			AliGraph	TLE	-	-	-	-	-	-
			Sancus	54.76	1.84	350.83	6.33	89.50	3.31	23086.58
		GCN-Large	Lotan	75.89	178.11	6.41	48.77	1.54	5917.79	3876.79
			Sancus	Fail	-	-	-	-	-	-
		GIN	Lotan	75.75	104.68	9.43	47.30	1.89	6735.52	4832.47
			DistDGL	Fail	-	-	-	-	-	-
			AliGraph	Fail	-	-	-	-	-	-
			Fail	-	-	-	-	-	-	-
#Nodes: 8.6M #Edges: 243.9M Avg. Degree: 28.2	amazon	GCN	Lotan	82.22	50.99	4.56	44.16	1.18	5451.33	3428.30
			DistDGL	86.14	261.58	0.02	10.66	51.19	4.48	4863.49
			AliGraph	Fail	-	-	-	-	-	-
			Sancus	Fail	-	-	-	-	-	-
		GIN	Lotan	91.79	252.84	2.26	39.37	1.08	4934.63	2230.87
			DistDGL	Fail	-	-	-	-	-	-
			AliGraph	Fail	-	-	-	-	-	-
			Fail	-	-	-	-	-	-	-
#Nodes: 111.1M #Edges: 1.6B Avg. Degree: 29.1	ogbn-papers100M	GCN	†Lotan	†-	†-	0.08	25.03	0.22	2499.20	801.93
			DistDGL	Fail	-	-	-	-	-	-
			AliGraph	Fail	-	-	-	-	-	-
			Sancus	Fail	-	-	-	-	-	-
		GIN	†Lotan	†-	†-	0.04	24.16	0.11	2530.61	811.36
			DistDGL	Fail	-	-	-	-	-	-
			AliGraph	Fail	-	-	-	-	-	-
			Fail	-	-	-	-	-	-	-

in their published papers. Table 2 first column shows brief statistics about the datasets.

Workloads. We define a GNN training workload with hyperparameter tuning factored in, which is an inevitable part of the end-to-end development of a GNN model. We primarily focus on two model architectures: GCN [26] and GIN [69] with various hyperparameter configurations. DistDGL and AliGraph have the batch size to tune additionally. In the corresponding literature, we found a batch size from 128 to 8192 is common. We tried as much as possible to make the comparison apples-to-apples and tune the batch size beforehand for them. To not understate their performance, we set the batch size to be as large as they could handle before failing to enable the maximum possible throughput. This means mini-batch size 8 for DistDGL on Amazon, 128 for DistDGL on ogbn-products+GCN and 8192 on ogbn-arxiv+GIN. And mini-batch size 128 for AliGraph on ogbn-arxiv+GCN. For Sancus, we can only test it on the GCN workloads as it does not have an existing implementation for GIN.

Experiment Setup. We use one cluster on CloudLab [52] with 8 worker nodes. Each node has two Intel Xeon 10-core 2.20 GHz CPUs, 192GB memory, and 10 Gbps network. Each worker node also has an Nvidia P100 GPU, which has 12 GB memory. We tried

to get GPUs with larger memory but such resources are scarce and costly to obtain, especially for the long-running tests we do. Nevertheless, it is not a showstopper as Lotan’s scalability gain is agnostic to the underlying hardware. Furthermore, even with larger GPUs, workloads can still scale beyond GPU memory capacity and would not change our observations about Lotan’s scalability. All nodes run Ubuntu 20.04. We use Spark 3.2.0, Pytorch 1.10, and CUDA 11.0.

7.1 End-to-end Performance Study

We use a 3-layer GCN with a hidden layer size of 256, as described in [21], dubbed GCN. We also include a variant of it with hidden size 512, which we call GCN-Large, to further distinguish between Lotan and Sancus. We skipped DistDGL and AliGraph with GCN-Large due to their crashes or much longer runtimes, and these tests would not provide extra insights. For GIN, we use one from [69] that is 4-layer. For the MLPs in GIN, we use a 2-layer MLP with dimensions {128, 256} for the end-to-end study. For the GCNs, their ApplyVertex functions are single-layer perceptions, while the GIN model uses the MLP described above. All of these models do not employ an ApplyEdge function and use a summation as the Aggregation.

Following the standard practices [21, 26, 69], we use an early termination of 10 epochs based on the validation set; we terminate if the validation accuracy does not increase for 10 consecutive epochs (with a tolerance of 0.01%). Further, we put a hard time limit of 48 hrs for each model config. We also combine the hyper-parameters used in the papers above to form a grid search: learning rate in $\{0.05, 0.01\}$, optimizer in $\{Adam, Adagrad\}$, and dropout in $\{0, 0.5\}$.

Table 2 summarizes the results of our end-to-end tests. On the ogbn-products + GCN workload, Lotan achieves 47x higher throughput than DistDGL while providing the same level of accuracy. There is no consensus from the GNN model research community on whether full-batch or mini-batch training is superior. Further, note some of the models and systems adopted mini-batch training partly due to the scalability issues of full-batch training [6, 18]. Lotan is designed to mitigate the said issues. Our finding of full-batch training achieving the same or slightly higher accuracy than mini-batch training is in line with prior work [24, 61]. Sancus, though it runs fast, has severe issues in accuracy during our test, likely due to its approximate nature. Furthermore, it starts to fail on the GCN-Large workload, but Lotan can still scale. Lotan is also the only system to be able to handle GIN training, and all other systems fail due to GPU memory issues. Increasing GPU memory might fix their problems on these specific workloads. Still, it would not resolve the fundamental issues these systems have and would not change the argument that Lotan has better scalability to handle large workloads.

On the tiny ogbn-arxiv dataset, while Lotan can still provide the highest accuracy on both GNNs, it no longer offers higher throughput than DistDGL. On the reddit dataset, which, despite having a similar number of nodes to ogbn-arxiv, has more edges, both DistDGL and AliGraph fail, likely due to the density of the graph. Sancus is still capable of operating and appears not affected much, but as other experiments showed, it offers lower accuracy due to its approximate nature. On the amazon dataset, Lotan and DistDGL are the only systems able to run the GCN workloads, and only Lotan for the GIN workloads. Lotan can provide a higher throughput than DistDGL. On the ogbn-papers100M dataset, one of the largest benchmark datasets available, Lotan is the only system able to run the workload. As far as we know, this is the first time for a GNN system to demonstrate full-graph GCN with a hidden size as large as 256 on this dataset. However, the execution is heavily bottlenecked, and a huge amount of disk spills happen. Consequently, we could not run the workload to converge in any reasonable amount of time. We only report the throughput numbers.

Figure 9 shows the learning curves for the best model out of some of the hyperparameter tuning workloads. On all workloads, Lotan converges fast and reaches the same level of accuracy as the SOTA. Regarding resource utilization, Lotan has high CPU utilization but generally lower GPU utilization across the workloads. This is because Lotan puts neural network operations on GPU and graph operations on CPU, while other systems put both on GPU. Except for Lotan, all other systems showed little to no disk R/W because they are not secondary-storage-aware, whereas Lotan can utilize the disk for spilling. Sancus and DistDGL further utilize GPU for communications, resulting in seemingly higher GPU utilization.




Figure 11: (A) Runtime breakdowns. (B) Ablation study.

7.2 Drill-down Experiments

To dig into the runtime figures, we also break down Lotan’s runtime and investigate each portion’s time costs in Figure 11(A). The Graph Engine costs dominate, especially on the larger dataset. DL Engine and Pipe-Join costs are not as significant. This composition will change when we try scaling the model in Section 7.2.2.

7.2.1 Ablation Study.

To inspect each component’s performance, we conduct an ablation study where we add our innovations to a naively implemented version of Lotan. We pick the ogbn-arxiv+GCN workload for this test. Figure 11(B) shows the results. We separate our technical innovations into four modules: (1) Reverse Graph Backprop (RGB) and the coupled GNN-centric partitioning scheme, as described in Section 5.1. (2) The execution plan rewrites by our Planner outlined in Section 4.2. (3) The various efforts we put into optimizing our Messenger architecture as described in Section 4.3. (4) Finally, our GNN Model Batching scheme proposed in Section 5.2.

All of the components have substantial contributions to performance gains; Reverse Graph Backprop can boost the performance by 2x without any plan rewrites or other optimizations. With Planner rewrites introduced, we get another 5x speed-up due to the sheer amount of communication and computation saved. Furthermore, our Messenger optimizations boost the performance by another 40% by reducing overheads in I/O, IPC, and synchronization. Last but not least, GNN Model Batching contributes a more than 5x speed-up due to amortized graph data access overheads. Overall, our technical innovations can boost the throughput of GNN model training by 76x, compared to a naively implemented system.

7.2.2 Model Scalability.

We now test Lotan’s capability of scaling to larger neural network models. In practice, there are two primary ways to scale up a model: make it deeper by adding more (GNN) layers, or increase the number of neurons in each layer. We call the first type depth scaling and the latter type width scaling. Since Lotan disaggregates the graph operations from neural network operations, it has very different behavior for the two types of scaling. To thoroughly test it, we use two different workloads based on the GIN model used earlier and train them on the ogbn-products dataset. For the depth scaling test, we test with different numbers of GNN layers ranging from 4 to 16. For the width scaling test, we fixed the model to be a 4-layer, and




Figure 12: Depth Scaling. (A) Runtime. (B) Utilization.




Figure 13: Width Scaling. (A) Runtime. (B) Utilization.

we varied the size of the MLP in GIN from 128 to 2^{17} (131072); we kept the embedding size also fixed as 256. This results in various models with hugely different sizes.

Depth Scaling. Figure 12 shows the results: Lotan can easily achieve almost linear scaling to 16 layers and even beyond, and there is minimal fluctuation in the processor utilizations. Note that the scaling is linear but not proportional; when the number of layers doubles, the runtime does not; this is because the scaling follows $y = kx + b$ with a non-zero intercept. It is to be expected as the amount of work grows linearly in this case, and Lotan shows resilience at scale. To the best of our knowledge, Lotan is the first system to demonstrate scale to 10+ layer GNNs with full batch training. It is important to note that the systems we compared all failed at 4 or more layers, as already discussed in Section 7.1.

Width Scaling. We show the width scaling results in Figure 13. Increasing the MLP size will not increase the amount of work on the Graph Engine, and since the GPU was under-utilized when the NN is small, we see an almost constant scaling of Lotan. In this case, we see a dramatic increase in GPU utilization and almost constant CPU utilization. Thanks to the decoupling of graph and neural networks, scaling one side does not necessarily affect the other. Lotan can provide independent scaling and frees the user from scalability issues. It enables the user to design the GNN components separately and more freely. Furthermore, Lotan can gracefully handle a GNN model with 140M+ parameters with full batch training. To put it into perspective, this is the number of parameters of some early Transformer DL models have: BERT (110M) [10], and GPT-1

(117M) [51]. To our knowledge, Lotan is the first system to be able to handle this scale among the GNN systems. As shown in Section 7.1, other systems all failed at the very beginning.

7.2.3 Model Batching.

We now inspect the effect of model batching on the workloads. For this test, we take the same ogbn-arxiv+GCN models used in Section 7.1 and create workloads with various degrees of model batching. Figure 14 shows the results. We first notice that the time costs scaling is all linear with constant overheads (manifested as the intercept), per our cost model described in Section 6. There is also a substantial gain in throughput, especially at the low degree of the model batching regime. The SGC and AAA costs scale far less steeply than the SGC costs; therefore, as the model batching size increases, the SGC costs become more and more dominant. This indicates that the biggest challenge is on graph data processing.

At a low degree of model batching (< 10), the time costs are dominated by their respective constant parts and not scaling as much with the model batch size. Therefore, the time costs only increase around 3x while the model batching size rises from 1 to 10, resulting in throughput gains. However, as the degree of model batching increases, the scaling parts of time costs dominate, and we see 2x increase in time costs when batch size increases from 10 to 20 (2x increase). Consequently, the throughput scaling plateaus out as 2x model batched would mean 2x more runtime in this realm.

8 RELATED WORK

GNN Systems. Many systems have been proposed to tackle the efficiency and scalability challenges of GNN training. Our work differs from them in our fundamental architecture design of separation of graph and neural network and our technical innovations. We have also conceptually compared them in Section 2.3 and tested against some of the most related and state-of-art systems in Section 7. Most of their techniques are complementary to our work. DGL [62, 75], and PyG [12] are prominent examples of all-purpose GNN frameworks designed for generality and usability. AliGraph [76], GraphScope [68], and PSGraph [25] are GNN systems designed for industry-scale usage with an emphasis on sampling-based GNN training, which differs from Lotan’s focus on full-batch training. NeuGraph [39] is one of the first systems to incorporate GNNs into an extended Gather-Apply-Scatter framework. It provides a scheduling scheme for shipping models/data in and out of multiple GPUs; its techniques are largely complementary to our work.

Other GNN systems proposed techniques ranging from memory management, communication reduction, approximated processing, and disk spilling. PaGraph [35] utilizes spare GPU memory for data caching to boost speed when the workload is relatively small. P3 [13] separates the graph metadata and graph properties and places them in a way to reduce communications. Sancus [50] proposes a communication reduction scheme via historical gradient caching and update skipping. Similarly, PipeGCN [61] uses pipeline parallelism with stale updates to speed up GNN training. They largely focus on the efficiency of GNN training via approximated processing and assumes the model and data can comfortably fit in GPU memory. PaGraph, P3, and Sancus are largely orthogonal to our work as Lotan is designed for large workloads, we do




Figure 14: Scaling with GNN Model Batching. (A) Throughput. (B) Time Costs. (C) Disk and Network Usage. (D) Utilization.

not assume an abundance of GPU memory. Dorylus [58] employs serverless functions to explore monetary cost-efficiency; our system is still for provisioned clusters, and we rely on existing data systems instead of custom-built ones used in Dorylus. Roc [24] uses main memory as swapping space to offload over-the-size data from GPU. MariusGNN [60] further proposes disk-spilling to increase the effective memory size. These techniques complement Lotan, and by employing a secondary-storage-aware graph data system, Lotan can naturally piggyback on its disk spilling capability. ALG [71] is designed for active learning setup which is largely orthogonal to our work. G3 [36] proposes to substitute DL frameworks with GPU-based graph operations; it can potentially be a candidate for the DL Engine in Lotan and is complementary to our work.

Graph Analytics Systems. Prior to the GNN era, large-scale systems were built for non-GNN graph analytical workloads and data management. Ranging from graph DBMS [9, 11, 29, 49], to classical graph analytics systems [2, 15–17, 27, 40, 57], and Graph Embedding learning (not to be confused with GNN) systems [30, 42, 65]. They are generally orthogonal to our work, as they target very different sets of workloads, and they seldom work with GNNs. Most techniques are workload-specific and not directly applicable to GNN training, but some may be complementary.

Faster and More Scalable GNNs. Ever since the first wave of GNN models arrived, algorithmic research has been active in tackling some of the scalability issues of GNNs by approximated processing and simplified architectures. This line of research is orthogonal to our work, as our goal is not to propose any new GNN model architecture but instead focus on the fundamental system challenges that will not be fixed by GNN model research alone. GraphSage [18] proposes mini-batch training and neighborhood sampling to reduce the data dependencies. FastGCN [6] runs even more aggressively IID sampling on the graph by directly controlling the number of nodes involved. SGC [63] challenges GNNs by proposing trivial two-layer architectures that reportedly could offer a similar accuracy performance. EIGNN [37] further extends SGC to an infinite depth model and uses eigendecomposition to boost efficiency. Graph coarsening techniques [22] have also been explored to preprocess and down-sample the input data.

9 LIMITATIONS, DISCUSSION, AND FUTURE WORK

Conclusions and limitations. By carefully abstracting, optimizing, and testing, we have demonstrated that it is possible to bridge the gap between graph analytics systems and DL systems with high scalability and without modifying their internal code. Currently, Lotan has two major limitations: (1) Lotan currently is only optimized for full-batch training. Mini-batch training would require efficient graph sampling and filtering, posing another scalability challenge and potential query optimization questions. (2) Lotan is meant for large workloads with large graphs and/or models. There is still some room for improvement on smaller workloads that do fit in memory, where more leeway for caching and batching techniques exists. Furthermore, GNN algorithmic research has accelerated in recent years, and more sophisticated architectures, some with over a thousand layers and some with Large Language Models (LLMs) baked in have been studied [31, 74]. More work is needed to see how Lotan or other GNN systems would fare for those models. We leave it to future work to mitigate these limitations.

Discussion. Our results showed that the graph data system bottlenecked many of our tests (see Figure 11(A)). There are many sync barriers, costly data replications, and frequent garbage collections. Graph data systems need to evolve to better support GNN workloads and property-rich graphs with high dimensional dense vectors. Furthermore, GNN systems such as Lotan can be extended to adopt recent advances in ML systems research for optimizing model selection workloads [28, 33, 45, 46, 73], fine-tuning and transfer learning workloads [44], and large model scaling [43, 47]. It is non-trivial to extend their techniques designed for IID data to graph data. However, with suitable adaptation, they could further amortize some runtime overheads to make GNNs more efficient at scale.

ACKNOWLEDGMENTS

This work was supported in part by an NSF CAREER Award under award number 1942724 and a gift from VMware. The content is solely the responsibility of the authors and does not necessarily represent the views of any of these organizations. We thank the members of UC San Diego’s Database Lab and Center for Networked Systems for their feedback on this work.

REFERENCES

- [1] M. Abadi, P. Barham, J. Chen, Z. Chen, A. Davis, J. Dean, M. Devin, S. Ghemawat, G. Irving, M. Isard, M. Kudlur, J. Levenberg, R. Monga, S. Moore, D. G. Murray, B. Steiner, P. A. Tucker, V. Vasudevan, P. Warden, M. Wicke, Y. Yu, and X. Zheng. Tensorflow: A system for large-scale machine learning. In *OSDI*, pages 265–283. USENIX Association, 2016.
- [2] S. Aridhi, A. Montresor, and Y. Velegrakis. BLADYG: A graph processing framework for large dynamic graphs. *Big Data Res.*, 9:9–17, 2017.
- [3] Y. Bai, C. Li, Z. Lin, Y. Wu, Y. Miao, Y. Liu, and Y. Xu. Efficient data loader for fast sampling-based gnn training on large graphs. *IEEE Transactions on Parallel & Distributed Systems*, (01):1–1, 2021.
- [4] M. Besta, E. Peter, R. Gerstenberger, M. Fischer, M. Podstawski, C. Barthels, G. Alonso, and T. Hoefler. Demystifying graph databases: Analysis and taxonomy of data organization, system designs, and graph queries. *CoRR*, abs/1910.09017, 2019.
- [5] A. Buluç, H. Meyerhenke, I. Safro, P. Sanders, and C. Schulz. Recent advances in graph partitioning. In *Algorithm Engineering*, volume 9220 of *Lecture Notes in Computer Science*, pages 117–158. 2016.
- [6] J. Chen, T. Ma, and C. Xiao. Fastgcn: Fast learning with graph convolutional networks via importance sampling. In *ICLR (Poster)*. OpenReview.net, 2018.
- [7] J. Chen, J. Zhu, and L. Song. Stochastic training of graph convolutional networks with variance reduction. In *ICML*, volume 80 of *Proceedings of Machine Learning Research*, pages 941–949. PMLR, 2018.
- [8] M. Chen, Z. Wei, Z. Huang, B. Ding, and Y. Li. Simple and deep graph convolutional networks. In *ICML*, volume 119 of *Proceedings of Machine Learning Research*, pages 1725–1735. PMLR, 2020.
- [9] A. Deutsch, Y. Xu, M. Wu, and V. E. Lee. Tigergraph: A native MPP graph database. *CoRR*, abs/1901.08248, 2019.
- [10] J. Devlin, M. Chang, K. Lee, and K. Toutanova. BERT: pre-training of deep bidirectional transformers for language understanding. In *NAACL-HLT (1)*, pages 4171–4186. Association for Computational Linguistics, 2019.
- [11] X. Feng, G. Jin, Z. Chen, C. Liu, and S. Salihoglu. Kuzu graph database management system. In *CIDR*, 2023.
- [12] M. Fey and J. E. Lenssen. Fast graph representation learning with pytorch geometric. *CoRR*, abs/1903.02428, 2019.
- [13] S. Gandhi and A. P. Iyer. P3: distributed deep graph learning at scale. In *OSDI*, pages 551–568. USENIX Association, 2021.
- [14] J. Gilmer, S. S. Schoenholz, P. F. Riley, O. Vinyals, and G. E. Dahl. Neural message passing for quantum chemistry. In *ICML*, volume 70 of *Proceedings of Machine Learning Research*, pages 1263–1272. PMLR, 2017.
- [15] Giraph. Apache Giraph. Accessed Nov 2, 2022. <https://giraph.apache.org/>.
- [16] J. E. Gonzalez, Y. Low, H. Gu, D. Bickson, and C. Guestrin. Powergraph: Distributed graph-parallel computation on natural graphs. In *OSDI*, pages 17–30. USENIX Association, 2012.
- [17] J. E. Gonzalez, R. S. Xin, A. Dave, D. Crankshaw, M. J. Franklin, and I. Stoica. Graphx: Graph processing in a distributed dataflow framework. In *OSDI*, pages 599–613. USENIX Association, 2014.
- [18] W. L. Hamilton, Z. Ying, and J. Leskovec. Inductive representation learning on large graphs. In *NIPS*, pages 1024–1034, 2017.
- [19] R. He and J. J. McAuley. Ups and downs: Modeling the visual evolution of fashion trends with one-class collaborative filtering. In J. Bourdeau, J. Hendler, R. Nkambo, I. Horrocks, and B. Y. Zhao, editors, *Proceedings of the 25th International Conference on World Wide Web, WWW 2016, Montreal, Canada, April 11 - 15, 2016*, pages 507–517. ACM, 2016.
- [20] L. Hoang, X. Chen, H. Lee, R. Dathathri, G. Gill, and K. Pingali. Efficient distribution for deep learning on large graphs. In *MLSys GNNsSys Workshop*. mlsys.org, 2021.
- [21] W. Hu, M. Fey, M. Zitnik, Y. Dong, H. Ren, B. Liu, M. Catasta, and J. Leskovec. Open graph benchmark: Datasets for machine learning on graphs. In *NeurIPS*, 2020.
- [22] Z. Huang, S. Zhang, C. Xi, T. Liu, and M. Zhou. Scaling up graph neural networks via graph coarsening. In *KDD*, pages 675–684. ACM, 2021.
- [23] A. Jain, A. R. Zamir, S. Savarese, and A. Saxena. Structural-rnn: Deep learning on spatio-temporal graphs. In *CVPR*, pages 5308–5317. IEEE Computer Society, 2016.
- [24] Z. Jia, S. Lin, M. Gao, M. Zaharia, and A. Aiken. Improving the accuracy, scalability, and performance of graph neural networks with roc. In I. S. Dhillon, D. S. Papaliopoulos, and V. Sze, editors, *Proceedings of Machine Learning and Systems 2020, MLSys 2020, Austin, TX, USA, March 2-4, 2020*. mlsys.org, 2020.
- [25] J. Jiang, P. Xiao, L. Yu, X. Li, J. Cheng, X. Miao, Z. Zhang, and B. Cui. Psgraph: How tencent trains extremely large-scale graphs with spark? In *ICDE*, pages 1549–1557. IEEE, 2020.
- [26] T. N. Kipf and M. Welling. Semi-supervised classification with graph convolutional networks. In *ICLR (Poster)*. OpenReview.net, 2017.
- [27] S. Ko and W. Han. Turbograph++: A scalable and fast graph analytics system. In *SIGMOD Conference*, pages 395–410. ACM, 2018.
- [28] A. Kumar, S. Nakandala, Y. Zhang, S. Li, A. Gemawat, and K. Nagrecha. Cerebro: A Layered Data Platform for Scalable Deep Learning. In *CIDR*. www.cidrdb.org, 2021.
- [29] A. G. Labouseur, J. Birnbaum, P. W. Olsen, S. R. Spillane, J. Vijayan, J. Hwang, and W. Han. The g* graph database: efficiently managing large distributed dynamic graphs. *Distributed Parallel Databases*, 33(4):479–514, 2015.
- [30] A. Lerer, L. Wu, J. Shen, T. Lacroix, L. Wehrstedt, A. Bose, and A. Peysakhovich. Pytorch-biggraph: A large scale graph embedding system. In *MLSys*. mlsys.org, 2019.
- [31] G. Li, M. Müller, B. Ghanem, and V. Koltun. Training graph neural networks with 1000 layers. In M. Meila and T. Zhang, editors, *Proceedings of the 38th International Conference on Machine Learning, ICML 2021, 18-24 July 2021, Virtual Event*, volume 139 of *Proceedings of Machine Learning Research*, pages 6437–6449. PMLR, 2021.
- [32] M. Li, D. G. Andersen, J. W. Park, A. J. Smola, A. Ahmed, V. Josifovski, J. Long, E. J. Shekita, and B.-Y. Su. Scaling Distributed Machine Learning with the Parameter Server. In *OSDI*, 2014.
- [33] S. Li and A. Kumar. Towards an optimized GROUP BY abstraction for large-scale machine learning. *Proc. VLDB Endow.*, 14(11):2327–2340, 2021.
- [34] X. Li, X. Liu, L. Lu, X. Hua, Y. Chi, and K. Xia. Multiphysical graph neural network (MP-GNN) for COVID-19 drug design. *Briefings Bioinform.*, 23(4), 2022.
- [35] Z. Lin, C. Li, Y. Miao, Y. Liu, and Y. Xu. Pagraph: Scaling GNN training on large graphs via computation-aware caching. In R. Fonseca, C. Delimitrou, and B. C. Ooi, editors, *SoCC '20: ACM Symposium on Cloud Computing, Virtual Event, USA, October 19-21, 2020*, pages 401–415. ACM, 2020.
- [36] H. Liu, S. Lu, X. Chen, and B. He. G3: when graph neural networks meet parallel graph processing systems on gpus. *Proc. VLDB Endow.*, 13(12):2813–2816, 2020.
- [37] J. Liu, K. Kawaguchi, B. Hooi, Y. Wang, and X. Xiao. EIGNN: efficient infinite-depth graph neural networks. In *NeurIPS*, pages 18762–18773, 2021.
- [38] Y. Low, J. Gonzalez, A. Kyrola, D. Bickson, C. Guestrin, and J. M. Hellerstein. Graphlab: A new framework for parallel machine learning. In *UAI*, pages 340–349. AUAI Press, 2010.
- [39] L. Ma, Z. Yang, Y. Miao, J. Xue, M. Wu, L. Zhou, and Y. Dai. Neugraph: Parallel deep neural network computation on large graphs. In *USENIX Annual Technical Conference*, pages 443–458. USENIX Association, 2019.
- [40] G. Malewicz, M. H. Austern, A. J. C. Bik, J. C. Dehnert, I. Horn, N. Leiser, and G. Czajkowski. Pregel: a system for large-scale graph processing. In *SIGMOD Conference*, pages 135–146. ACM, 2010.
- [41] X. Miao, W. Zhang, Y. Shao, B. Cui, L. Chen, C. Zhang, and J. Jiang. Lasagne: A multi-layer graph convolutional network framework via node-aware deep architecture (extended abstract). In *ICDE*, pages 1561–1562. IEEE, 2022.
- [42] J. Mohoney, R. Waleffe, H. Xu, T. Rekatsinas, and S. Venkataraman. Marius: Learning massive graph embeddings on a single machine. In *OSDI*, pages 533–549. USENIX Association, 2021.
- [43] K. Nagrecha and A. Kumar. Saturn: An Optimized Data System for Multi-Large-Model Deep Learning Workloads. https://adalabucsd.github.io/papers/TR_2023_Saturn.pdf, 2023. [Tech report].
- [44] S. Nakandala and A. Kumar. Nautilus: An optimized system for deep transfer learning over evolving training datasets. In Z. G. Ives, A. Bonifati, and A. E. Abbadi, editors, *SIGMOD '22: International Conference on Management of Data, Philadelphia, PA, USA, June 12 - 17, 2022*, pages 506–520. ACM, 2022.
- [45] S. Nakandala, Y. Zhang, and A. Kumar. Cerebro: Efficient and Reproducible Model Selection on Deep Learning Systems. In *Proceedings of the 3rd International Workshop on Data Management for End-to-End Machine Learning*, pages 1–4, 2019.
- [46] S. Nakandala, Y. Zhang, and A. Kumar. Cerebro: A Data System for Optimized Deep Learning Model Selection. *Proc. VLDB Endow.*, 13(11):2159–2173, 2020.
- [47] D. Narayanan, A. Harlap, A. Phanishayee, V. Seshadri, N. R. Devanur, G. R. Ganger, P. B. Gibbons, and M. Zaharia. Pipedream: generalized pipeline parallelism for DNN training. In *SOSP*, pages 1–15. ACM, 2019.
- [48] D. Narayanan, K. Santhanam, A. Phanishayee, and M. Zaharia. Accelerating deep learning workloads through efficient multi-model execution. In *NeurIPS Workshop on Systems for Machine Learning*, December 2018.
- [49] Neo4j. Neo4j. Accessed October 12, 2021. <https://neo4j.com/>.
- [50] J. Peng, Z. Chen, Y. Shao, Y. Shen, L. Chen, and J. Cao. SANCUS: staleness-aware communication-avoiding full-graph decentralized training in large-scale graph neural networks. *Proc. VLDB Endow.*, 15(9):1937–1950, 2022.
- [51] A. Radford and K. Narasimhan. Improving language understanding by generative pre-training. 2018.
- [52] R. Ricci, E. Eide, and CloudLabTeam. Introducing Cloudlab: Scientific Infrastructure for Advancing Cloud Architectures and Applications. ; *login:: the magazine of USENIX & SAGE*, 39(6):36–38, 2014.
- [53] A. Roy, I. Mihailovic, and W. Zwaenepoel. X-stream: edge-centric graph processing using streaming partitions. In *SOSP*, pages 472–488. ACM, 2013.
- [54] S. Sakr, A. Bonifati, H. Voigt, A. Iosup, K. Ammar, R. Angles, W. G. Aref, M. Arenas, M. Besta, P. A. Boncz, K. Daudjee, E. D. Valle, S. Dumbrava, O. Hartig, B. Hashhofer, T. Hegeman, J. Hidders, K. Hose, A. Iamnitchi, V. Kalavri, H. Kapp, W. Martens, M. T. Özsu, E. Peukert, S. Plantikow, M. Ragab, M. Ripeanu, S. Salihoglu, C. Schulz, P. Selmer, J. F. Sequeda, J. Shinavier, G. Szárnyas, R. Tommasini, A. Tumeo, A. Uta, A. L. Varbanescu, H. Wu, N. Yakovets, D. Yan, and E. Yoneki. The future is

- big graphs: a community view on graph processing systems. *Commun. ACM*, 64(9):62–71, 2021.
- [55] A. Sergeev and M. D. Balso. Horovod: Fast and Easy Distributed Deep Learning in TF. *arXiv preprint arXiv:1802.05799*, 2018.
- [56] J. Sun, M. Yue, Z. Lin, X. Yang, L. Nocera, G. Kahn, and C. Shahabi. Crime-forecaster: Crime prediction by exploiting the geographical neighborhoods’ spatiotemporal dependencies. In *ECML/PKDD (5)*, volume 12461 of *Lecture Notes in Computer Science*, pages 52–67. Springer, 2020.
- [57] C. H. C. Teixeira, A. J. Fonseca, M. Serafini, G. Siganos, M. J. Zaki, and A. Aboul-naga. Arabesque: a system for distributed graph mining. In *SOSP*, pages 425–440. ACM, 2015.
- [58] J. Thorpe, Y. Qiao, J. Elyofson, S. Teng, G. Hu, Z. Jia, J. Wei, K. Vora, R. Netravali, M. Kim, and G. H. Xu. Dorylus: Affordable, scalable, and accurate GNN training with distributed CPU servers and serverless threads. In *OSDI*, pages 495–514. USENIX Association, 2021.
- [59] Y. Tian. The world of graph databases from an industry perspective. *CoRR*, abs/2211.13170, 2022.
- [60] R. Waleffe, J. Mohoney, T. Rekatsinas, and S. Venkataraman. Mariusgnn: Resource-efficient out-of-core training of graph neural networks, 2022.
- [61] C. Wan, Y. Li, C. R. Wolfe, A. Kyriallidis, N. S. Kim, and Y. Lin. Pipegcn: Efficient full-graph training of graph convolutional networks with pipelined feature communication. In *ICLR*. OpenReview.net, 2022.
- [62] M. Wang, L. Yu, D. Zheng, Q. Gan, Y. Gai, Z. Ye, M. Li, J. Zhou, Q. Huang, C. Ma, Z. Huang, Q. Guo, H. Zhang, H. Lin, J. Zhao, J. Li, A. J. Smola, and Z. Zhang. Deep graph library: Towards efficient and scalable deep learning on graphs. *CoRR*, abs/1909.01315, 2019.
- [63] F. Wu, A. H. S. Jr., T. Zhang, C. Fifty, T. Yu, and K. Q. Weinberger. Simplifying graph convolutional networks. In *ICML*, volume 97 of *Proceedings of Machine Learning Research*, pages 6861–6871. PMLR, 2019.
- [64] S. Wu, F. Sun, W. Zhang, X. Xie, and B. Cui. Graph neural networks in recommender systems: A survey. *ACM Comput. Surv.*, 55(5):97:1–97:37, 2023.
- [65] T. Wu, A. Khan, M. Yong, G. Qi, and M. Wang. Efficiently embedding dynamic knowledge graphs. *Knowl. Based Syst.*, 250:109124, 2022.
- [66] Z. Wu, S. Pan, F. Chen, G. Long, C. Zhang, and P. S. Yu. A comprehensive survey on graph neural networks. *IEEE Trans. Neural Networks Learn. Syst.*, 32(1):4–24, 2021.
- [67] H. Xie, D. Li, Y. Wang, and Y. Kawai. Visualization method for the spreading curve of COVID-19 in universities using GNN. In *BigComp*, pages 121–128. IEEE, 2022.
- [68] J. Xu, Z. Bai, W. Fan, L. Lai, X. Li, Z. Li, Z. Qian, L. Wang, Y. Wang, W. Yu, and J. Zhou. Graphscope: A one-stop large graph processing system. *Proc. VLDB Endow.*, 14(12):2703–2706, 2021.
- [69] K. Xu, W. Hu, J. Leskovec, and S. Jegelka. How powerful are graph neural networks? In *ICLR*. OpenReview.net, 2019.
- [70] R. Ying, R. He, K. Chen, P. Eksombatchai, W. L. Hamilton, and J. Leskovec. Graph convolutional neural networks for web-scale recommender systems. In *KDD*, pages 974–983. ACM, 2018.
- [71] W. Zhang, Y. Shen, Y. Li, L. Chen, Z. Yang, and B. Cui. ALG: fast and accurate active learning framework for graph convolutional networks. In *SIGMOD Conference*, pages 2366–2374. ACM, 2021.
- [72] Y. Zhang and A. Kumar. Lotan: Bridging the Gap between GNNs and Scalable Graph Analytics Engines. https://adalabucsd.github.io/papers/TR_2023_Lotan.pdf, 2023. [Tech report].
- [73] Y. Zhang, F. McQuillan, N. Jayaram, N. Kak, E. Khanna, O. Kislal, D. Valdano, and A. Kumar. Distributed deep learning on data systems: A comparative analysis of approaches. *Proc. VLDB Endow.*, 14(10):1769–1782, 2021.
- [74] J. Zhao, M. Qu, C. Li, H. Yan, Q. Liu, R. Li, X. Xie, and J. Tang. Learning on large-scale text-attributed graphs via variational inference. In *The Eleventh International Conference on Learning Representations*, 2023.
- [75] D. Zheng, C. Ma, M. Wang, J. Zhou, Q. Su, X. Song, Q. Gan, Z. Zhang, and G. Karypis. Distdgl: Distributed graph neural network training for billion-scale graphs. In *10th IEEE/ACM Workshop on Irregular Applications: Architectures and Algorithms, IA3 2020, Atlanta, GA, USA, November 11, 2020*, pages 36–44. IEEE, 2020.
- [76] R. Zhu, K. Zhao, H. Yang, W. Lin, C. Zhou, B. Ai, Y. Li, and J. Zhou. Aligraph: A comprehensive graph neural network platform. *Proc. VLDB Endow.*, 12(12):2094–2105, 2019.